

nf-core/viralmetagenome: A Novel Pipeline for Untargeted Viral Genome Reconstruction

Joon Klaps^{1,*}, Philippe Lemey¹, Magdalini Bletsas¹, nf-core community², Liana Eleni Kafetzopoulou^{1,3}

¹Rega Institute for Medical Research, Department of Microbiology, Immunology and Transplantation, KU Leuven, Belgium

²A full list of contributors can be found at <https://nf-co.re/community>.

³Leiden University Center of Infectious Diseases (LUCID), Leiden University Medical Center, Albinusdreef 2, 2333 ZA, Leiden, the Netherlands

*Correspondence: joon.klaps@kuleuven.be

November 30, 2025

Abstract

Motivation: ~~Eukaryotic viruses present significant challenges for genome reconstruction and variant analysis—Reconstructing eukaryotic viral genomes from metagenomic data is challenging~~ due to their extensive diversity and potential genome segmentation. ~~While de novo assembly followed by reference database matching and scaffolding is a commonly used approach, the manual execution of this workflow is extremely time-consuming, particularly due to the extensive reference curation required. Here, we~~ Current approaches often rely on labor-intensive manual curation for reference selection and scaffolding, limiting scalability for large studies or rapid outbreak response. We address the critical need for an automated, scalable pipeline ~~that can efficiently handle~~ for efficient viral metagenomic analysis without manual intervention.

Results: We present nf-core/viralmetagenome, a comprehensive ~~viral metagenomic pipeline for untargeted genome~~ Nextflow pipeline for the untargeted reconstruction and variant analysis of eukaryotic DNA and RNA viruses. ~~Viralmetagenome is implemented as a Nextflow workflow that processes from~~ short-read metagenomic samples to automatically detect and assemble viral genomes, while also performing variant analysis or hybridisation capture enriched samples. The pipeline ~~features~~ automates the entire process from read preprocessing to consensus generation, integrating multiple *de novo* assemblers, automated reference selection, ~~consensus quality control metrics, comprehensive~~ and iterative consensus refinement. It features robust quality control, extensive documentation, and seamless ~~integration with containerization technologies, including portability via Docker and Singularity.~~ We ~~demonstrate the utility and accuracy of our approach through validation on both~~ validated the pipeline on diverse simulated and real datasets, showing robust performance across diverse viral families in metagenomic samples demonstrating its ability to recover high-quality genomes from complex metagenomic samples and resolve co-infections, making it a powerful tool for viral surveillance.

Availability: nf-core/viralmetagenome is freely available at <https://github.com/nf-core/viralmetagenome> <https://github.com/nf-core/viralmetagenome> with comprehensive documentation at <https://nf-co.re/viralmetagenome>

co.re/viralmetagenome

Contact: joon.klaps@kuleuven.be

Supplementary information: Supplementary data are available at <https://github.com/Joon-Klaps/nf-core-viralmetagenome-manuscript> online.

Keywords: viralmetagenome, bioinformatic pipeline, nextflow, viral metagenomics, viral assembly, viral variant analysis

1 Introduction

Reconstructing viral genomes from metagenomic sequencing data presents considerable computational challenges, particularly for viruses exhibiting extensive genetic diversity (Baaijens et al. 2017, Meleshko et al. 2021). This ~~diversity is further compounded by segmented genomes in families like~~ challenge is amplified in segmented viruses such as influenza, rotavirus, and bunyaviruses, where individual segments can evolve under distinct selective pressures and ~~reassort, contributing to a complex landscape for genome reassortment~~ complicates reconstruction. While ~~pipelines are often designed to some~~ pipelines target specific viruses and their subtypes (Shepard et al. 2016), accurate ~~and complete genome reconstruction of genome reconstruction from~~ samples with unknown ~~references typically requires viral strains typically requires time-consuming~~ manual curation of contigs and reference matching (de Vries et al. 2021). This ~~manual curation process is time-consuming, making it impractical for makes~~ large-scale metagenome studies or rapid ~~response scenarios that involve emerging viral outbreaks of unknown origin~~ outbreak responses impractical.

To address these limitations, we developed nf-core/viralmetagenome, a comprehensive pipeline specifically designed for untargeted viral genome reconstruction ~~using short read sequences~~. The pipeline is developed using Nextflow (Di Tommaso et al. 2017) within the nf-core framework (Ewels et al. 2020)(Ewels et al. 2020, Langer et al. 2025), ensuring reproducibility through containerization with Docker (Merkel 2014) and Singularity (Kurtzer et al. 2017), and enabling portability across computational platforms such as local desktops, high-performance clusters and cloud environments.

2 Pipeline Description

nf-core/viralmetagenome implements an automated workflow ~~performing for de novo~~ assembly, reference matching, and iterative consensus refinement ~~for the reconstruction of to reconstruct~~ viral genomes without prior target knowledge. The pipeline ~~consists of comprises~~ five major stages: read preprocessing, metagenomic diversity assessment, contig assembly and scaffolding, iterative consensus refinement with variant analysis, and quality control (Figure 1). ~~While this manuscript highlights key differences between particular tools, unless~~ Unless otherwise specified, the pipeline offers multiple options to accommodate established user workflows and preferences. ~~In depth tool details are available~~ Detailed tool descriptions are provided in Supplementary Table 1.




Figure 1: ~~Visual overview of the~~ **Overview of the nf-core/viralmetagenome pipeline for untargeted viral genome reconstruction.** Default options are highlighted in bold. ~~nf-core/viralmetagenome pipeline for untargeted viral genome reconstruction.~~ ~~nf-core/viralmetagenome~~ processes short-read data through pre-processing, metagenomic diversity assessment, *de novo* assembly with multiple assemblers, scaffolding with automated reference identification and ~~contig~~-taxonomy-guided clustering of **contigs**, and iterative consensus refinement through read mapping and ~~variant~~-consensus calling. Quality control metrics, assembly statistics, and coverage data are integrated into interactive MultiQC reports and standardised overview tables for downstream analysis.

2.1 Read preprocessing

~~Input reads~~ Input reads are provided via sample sheets containing sample names and paths to short-read FASTQ ~~paths~~ files. These are preprocessed using FastQC and ~~have their adapters trimmed with trimmed for adapters using fastp~~ (Chen et al. 2018) (default) or Trimmomatic (Bolger et al. 2014). Fastp ~~is overall faster and has offers~~ automated adapter detection and ~~trimming faster processing~~ (Chen et al. 2018). UMI deduplication is implemented using HUMID (Laros 2025) ~~and once reads are mapped to a reference with for raw reads, or~~ UMI-tools after mapping (Smith et al. 2017). Multiple sequencing runs ~~are can be~~ merged after trimming by specifying merge group identifiers in the sample sheet. ~~Complexity~~ Optional filtering with bbdduk (Bushnell 2022) or PRINSEQ++ (Cantu et al. 2019) removes low-complexity sequences ~~containing repetitive elements where, while~~ host reads are removed with Kraken2 (Wood et al. 2019).

2.2 Metagenomic diversity assessment

Taxonomic classification of preprocessed reads is performed using two complementary ~~approaches—~~ ~~Kaiju~~ approaches—Kaiju (Menzel et al. 2016) and Kraken2 (Wood et al. 2019) ~~to—to~~ to maximise detection sensitivity across diverse viral families. Results from both classifiers are visualised using Krona (Ondov et al. 2011).

2.3 ~~De novo~~ De novo assembly and clustering

The assembly workflow ~~implements employs a~~ multi-assembler ~~approaches followed by clustering and scaffolding~~ approach where contigs generated by all tools are pooled and processed jointly. *De novo* assembly ~~is performed using uses~~ SPAdes (Meleshko et al. 2021) (RNAviral mode), MEGAHIT (Li et al. 2016), and Trinity (Grabherr et al. 2011), capitalizing on ~~their~~ distinct algorithmic strengths to maximise genome recovery across diverse viral families and variable read depths.

~~Reference identification uses BLASTn (Altschul et al. 1990)~~ Subsequently, a BLASTn (Altschul et al. 1990) search against the Reference Viral Database (RVDB) (Goodacre et al. 2018) ~~retaining top five hits to facilitate identification of related genomic segments and appropriate reference sequences for~~ identifies candidate references for each contig. All contigs and their respective references are then subjected to clustering. This strategy enables contig scaffolding against the cluster centroid—whether an external reference or a *de novo* contig ~~scaffolding and clustering~~—thereby combining outputs from multiple assemblers into unified scaffolds.

Clustering is performed in two sequential stages. First, taxonomic pre-clustering groups contigs based on taxonomic classification using Kraken2 (Wood et al. 2019) and Kaiju (Menzel et al. 2016), with optional filtering to focus on specific taxonomic clades for more targeted analyses. Second, nucleotide similarity clustering within taxonomic groups is performed using CD-HIT-EST (Li and Godzik 2006) (default), VSEARCH (Rognes et al. 2016), MMseqs2 (Steinegger and Söding 2017), vRhyme (Kieft et al. 2022), or Mash (Ondov et al. 2019). All tools are valid options, though performance may vary depending on the dataset (Zielezinski et al. 2025, Steinegger and Söding 2017).

~~As an optional filtering step of contig clusters, after~~ Optionally, following assembly and extension,

reads can be mapped to all contigs using BWA-MEM2 (Vasimuddin et al. 2019) (default), ~~BWA (?)~~, ~~or bowtie2~~ or Bowtie2 (Langmead et al. 2019). Clusters are filtered based on the total percentage of reads mapping to all contigs within a cluster, ~~allowing identification and removal of clusters that likely represent assembly artefacts resulting from low read coverage.~~ enabling the removal of assembly artefacts. For the final scaffolding step, all cluster members are mapped to the cluster representative or centroid using Minimap2 (Li 2018), followed by consensus calling with iVar (Grubaugh et al. 2019) to generate reference-assisted assemblies. ~~Regions with zero coverage depth can optionally be represented by the reference genome to produce a more complete scaffold genome for consensus calling.~~ scaffolds.

2.4 ~~Iterative consensus refinement~~ Consensus generation and variant calling

The consensus module supports external reference-based analysis and scaffold refinement. Users can provide ~~a separate reference genome or reference set for each sample with mapping constraints;~~ ~~when a reference~~ specific reference genomes; if a set is provided, Mash (Ondov et al. 2019) selects the most similar ~~can be selected using Mash (Ondov et al. 2019)~~ one. The selected reference genome undergoes a single round of refinement to generate the consensus.

~~Scaffold refinement performs~~ In contrast, scaffold refinement employs an iterative approach, performing up to 4 iterative cycles (default 2). ~~Each iteration maps to~~ polish the consensus sequence. Both processes map reads using BWA-MEM2 (Vasimuddin et al. 2019) (default), or Bowtie2 (Langmead et al. 2019), ~~BWA (?)~~, ~~or bowtie2 (Langmead et al. 2019)~~ to the consensus, followed by variant calling with BCFtools (Danecek et al. 2021) or iVar (Grubaugh et al. 2019). ~~Benchmarking by Bassano et al. (Bassano et al. 2022) showed that BCFtools outperformed (default). Benchmarking (Bassano et al. 2022) indicated that BCFtools outperforms iVar in precision and recall, where while iVar identified more low frequency low-frequency variants. Users are recommended to consider prioritising~~ can prioritise sensitivity or specificity when selecting the variant caller. Final variant calling is performed on the polished consensus to characterise single nucleotide variants.

2.5 Consensus Quality Control

Quality control employs CheckV (Nayfach et al. 2021) ~~for completeness estimates~~ to assess completeness, BLASTn (Altschul et al. 1990) for reference similarity, and MMseqs2 (Steinegger and Söding 2017) against the annotated database ViroSaurus (Gleizes et al. 2020). These analyses enable species identification, genomic segment classification, host prediction, and ~~any other retrieval~~ of additional metadata embedded within the reference databases.

The ~~refinement progression~~ consensus refinement is evaluated through sequence alignment with MAFFT (Katoh et al. 2002), which compares final consensus genomes against *de novo* contigs, intermediate consensus sequences from iterative cycles, and the scaffolding reference. All tool metrics are compiled into an interactive MultiQC report (Ewels et al. 2016). Additionally, key metrics are extracted from the MultiQC report and compiled into standalone overview tables to facilitate downstream analysis across all processed samples.

3 Applications

~~To assess the performance of nf-core/viralmetagenome under challenging scenarios, we simulated coinfection scenarios by mixing paired-end reads from public HIV-1 genomes with varying diversity (80-99%), resulting in 13 samples (See supplementary table 2). Nf-core/viralmetagenome successfully identified coinfections in all mixed samples when genetic similarity was low to moderate ($\leq 96.7\%$).~~

~~We validated nf-core/viralmetagenome performance on~~ We validated the pipeline's performance using a diverse set of real metagenomic samples ~~spanning~~ containing human and plant pathogens. ~~Here, the pipeline successfully generated nf-core/viralmetagenome successfully reconstructed~~ high-quality or near-complete ~~genomes across viral families including segmented consensus genomes across~~ various viral families, including segmented viruses (Lassa virus, ~~Orthonairovirus~~ Orthonairovirus, Tomato spotted wilt tospovirus) and non-segmented viruses (SARS-CoV-2, West Nile virus, Potato virus Y, Youcai mosaic virus, and Monkeypox virus).

~~Processing;~~ Supplementary methods S1). The pipeline proved efficient, processing 28 samples ~~(supplementary methods) required in~~ 412 CPU hours ~~and a maximum~~ with a maximum memory footprint of 79GB RAM ~~on an HPC, excluding taxonomic classification steps.~~ The automated reference selection offers substantial improvements over manual curation by reducing processing time while preserving reconstruction accuracy. Performance correlates strongly with reference database comprehensiveness, as consensus genomes tended to be more complete and similar to the true consensus sequence when the scaffolding reference was closer to the true viral genome on a high-performance cluster. To facilitate local deployment, we implemented a `-profile local` option, which utilizes a smaller viral RefSeq database instead of the full RVDB database for Kaiju, reducing the pipeline's peak RAM usage to approximately 18GB.

To evaluate the pipeline's robustness in complex scenarios, we conducted two targeted experiments. First, we assessed its ability to resolve co-infections using a simulated mixture of viral genomes (Supplementary Methods S2). The pipeline successfully distinguished and reconstructed distinct co-infecting strains, provided the genetic divergence was sufficient ($\leq 88.7\%$ ANI; Supplementary Figure 1).

The second part of our investigation explored the critical role of the scaffolding reference in consensus genome reconstruction. ~~This emphasizes~~ We selected sequencing data from seven Lassa virus-positive human clinical samples with diverse viral loads (Supplementary Methods S3; Supplementary Figure 2). For these seven samples, a significantly higher sequencing depth was observed for the S segment than for the L segment (Supplementary Figure 3). To establish a benchmark, a baseline consensus genome was first generated for each sample using automated reference selection from the full unclustered-RVDB (U-RVDB) reference pool. We then systematically challenged the scaffolding process by providing it a reference genome from a range of references that have varying degrees of similarity to the baseline consensus (64%–100% ANI; Supplementary Figure 4).

Comparing these generated consensus sequences against their respective baseline highlights the utility of reference scaffolding (Figure 2), particularly in low read depth scenarios when *de novo* assembly was challenging. For the high-coverage S segment, *de novo* assembly was consistently sufficient to reconstruct the full genome, rendering the choice of scaffolding reference largely redundant (Figure 2A-B). In contrast, the low-coverage L segment proved more challenging: *de novo* assem-

bly produced shorter, incomplete fragments, demonstrating that the pipeline’s ability to identify and utilize a closely related external reference was essential to bridge gaps and recover a complete segment consensus genome (Figure 2C-D).




Figure 2: **Impact of scaffolding reference similarity on consensus reconstruction accuracy.** The baseline consensus for each sample was generated using the pipeline’s automated reference selection from the full unclustered-RVDB (U-RVDB) pool. (A, B) Boxplots comparing Consensus Accuracy (% ANI) for the L and S segments. The plots contrast the performance of using an external sequence as the scaffolding backbone (green) against relying solely on the *de novo* contig (orange). For the low-coverage L segment (A), external references significantly improve accuracy, whereas for the high-coverage S segment (B), the *de novo* contig performs comparably to reference scaffolding (sequencing depths detailed in Supplementary Figure S3). (C, D) Scatter plots illustrating the relationship between reference similarity (% ANI) and consensus accuracy (% ANI). The dashed line represents the line of identity ($y = x$), and points are colored by sample source (as per Supplementary Figure S2).

These findings emphasise the need to keep databases like RVDB (Goodacre et al. 2018) and Virosaurus (Gleizes et al. 2020) up-to-date. Since nf-core/viralmetagenome is primarily designed for eukaryotic viruses, bacteriophage analysis requires different approaches and users interested in bacteriophage analysis are encouraged to explore pipelines targeting designed to target phages such as VIRify (Rangel-Pineros et al. 2022), VIBRANT (Kieft et al. 2020), VirSorter2 (Guo et al. 2021).

4 Conclusion

nf-core/viralmetagenome addresses a critical need in viral genomics by providing an automated, scalable solution for untargeted viral genome reconstruction. The pipeline successfully automates the traditionally time-consuming and manual ~~execution~~ process of viral genome ~~assembly-reconstruction~~ from short-read metagenomic data through its integrated workflow of *de novo* contig assembly, automated reference selection, clustering algorithms, and iterative refinement strategies.

Our validation ~~demonstrates the pipeline's broad applicability across diverse results~~ indicate that the pipeline is suitable for a wide range of eukaryotic viral families, ~~achieving-it~~ removes the burden of extensive manual curation, and enables consistent, high-quality genome reconstruction ~~while ensuring reproducibility and ease of deployment across different computational environments.~~

with reproducible results across varied computational settings. As viral surveillance and outbreak response increasingly rely on metagenomic sequencing, automated pipelines like nf-core/viralmetagenome will be essential for the timely identification of pathogen strains. The pipeline represents a significant step forward in making viral genome reconstruction accessible to researchers without requiring extensive bioinformatics expertise, facilitating broader adoption of metagenomic approaches in viral research and public health applications.

Acknowledgments

J.K, P.L. and L.E.K acknowledge support from the Research Foundation - Flanders (Fonds voor Wetenschappelijk Onderzoek – Vlaanderen, G005323N and G051322N, 1SH2V24N, 12X9222N). Assistance from AI language models (e.g., ChatGPT) was used to enhance the clarity of the manuscript and to support code development.

Author Contributions

J.K. designed and implemented the pipeline, performed validation analyses, and wrote the manuscript. P.L. and L. E. K. supervised the project and along with M.B. provided critical feedback. The nf-core community contributed to maintaining the pipeline. All authors reviewed and approved the final manuscript.

Conflict of Interest

The authors declare no competing interests.

References

S F Altschul, W Gish, W Miller, E W Myers, and D J Lipman. Basic local alignment search tool. *J. Mol. Biol.*, 215(3):403–410, 1990. doi: 10.1016/S0022-2836(05)80360-2.

- 232 Jasmijn A Baaijens, Amal Zine El Aabidine, Eric Rivals, and Alexander Schönhuth. De novo
233 assembly of viral quasiespecies using overlap graphs. *Genome Res.*, 27(5):835–848, 2017. doi:
234 10.1101/gr.215038.116.
- 235 Irene Bassano, Vinoy K Ramachandran, Mohammad S Khalifa, Chris J Lilley, Mathew R Brown,
236 Ronny van Aerle, Hubert Denise, William Rowe, Airey George, Edward Cairns, Claudia
237 Wierzbicki, Natalie D Pickwell, Myles Wilson, Matthew Carlile, Nadine Holmes, Alexander Payne,
238 Matthew Loose, Terry A Burke, Steve Paterson, Matthew J Wade, and Jasmine M S Grim-
239 sley. Evaluation of variant calling algorithms for wastewater-based epidemiology using mixed
240 populations of SARS-CoV-2 variants in synthetic and wastewater samples. *bioRxiv*, 2022. doi:
241 10.1101/2022.06.06.22275866.
- 242 Anthony M Bolger, Marc Lohse, and Bjoern Usadel. Trimmomatic: a flexible trimmer for illumina
243 sequence data. *Bioinformatics*, 30(15):2114–2120, 2014. doi: 10.1093/bioinformatics/btu170.
- 244 Brian Bushnell. BBMap, 2022.
- 245 Vito Adrian Cantu, Jeffrey Sadural, and Robert Edwards. PRINSEQ++, a multi-threaded tool
246 for fast and efficient quality control and preprocessing of sequencing datasets. *PeerJ*, 2019. doi:
247 10.7287/peerj.preprints.27553v1.
- 248 Shifu Chen, Yanqing Zhou, Yaru Chen, and Jia Gu. fastp: an ultra-fast all-in-one FASTQ prepro-
249 cessor. *Bioinformatics*, 34(17):i884–i890, 2018. doi: 10.1093/bioinformatics/bty560.
- 250 Petr Danecek, James K Bonfield, Jennifer Liddle, John Marshall, Valeriu Ohan, Martin O Pollard,
251 Andrew Whitwham, Thomas Keane, Shane A McCarthy, Robert M Davies, and Heng Li. Twelve
252 years of SAMtools and BCFtools. *Gigascience*, 10(2), 2021. doi: 10.1093/gigascience/giab008.
- 253 Jutte J C de Vries, Julianne R Brown, Nicole Fischer, Igor A Sidorov, Sofia Morfopoulou, Jiabin
254 Huang, Bas B Oude Munnink, Arzu Sayiner, Alihan Bulgurcu, Christophe Rodriguez, Guillaume
255 Gricourt, Els Keyaerts, Leen Beller, Claudia Bachofen, Jakub Kubacki, Cordey Samuel, Laubscher
256 Florian, Schmitz Dennis, Martin Beer, Dirk Hoeper, Michael Huber, Verena Kufner, Maryam
257 Zaheri, Aitana Lebrand, Anna Papa, Sander van Boheemen, Aloys C M Kroes, Judith Breuer,
258 F Xavier Lopez-Labrador, and Eric C J Claas. Benchmark of thirteen bioinformatic pipelines for
259 metagenomic virus diagnostics using datasets from clinical samples. *J. Clin. Virol.*, 141:104908,
260 2021. doi: 10.1016/j.jcv.2021.104908.
- 261 Paolo Di Tommaso, Maria Chatzou, Evan W Floden, Pablo Prieto Barja, Emilio Palumbo, and
262 Cedric Notredame. Nextflow enables reproducible computational workflows. *Nat. Biotechnol.*, 35
263 (4):316–319, 2017. doi: 10.1038/nbt.3820.
- 264 Philip Ewels, Måns Magnusson, Sverker Lundin, and Max Käller. MultiQC: summarize analysis
265 results for multiple tools and samples in a single report. *Bioinformatics*, 32(19):3047–3048, 2016.
266 doi: 10.1093/bioinformatics/btw354.
- 267 Philip A Ewels, Alexander Peltzer, Sven Fillinger, Harshil Patel, Johannes Alneberg, Andreas
268 Wilm, Maxime Ulysse Garcia, Paolo Di Tommaso, and Sven Nahnsen. The nf-core framework
269 for community-curated bioinformatics pipelines. *Nat. Biotechnol.*, 38(3):276–278, 2020. doi:
270 10.1038/s41587-020-0439-x.

- 271 Anne Gleizes, Florian Laubscher, Nicolas Guex, Christian Iseli, Thomas Junier, Samuel Cordey,
272 Jacques Fellay, Ioannis Xenarios, Laurent Kaiser, and Philippe Le Mercier. Virosaurus a reference
273 to explore and capture virus genetic diversity. *Viruses*, 12(11), 2020. doi: 10.3390/v12111248.
- 274 Norman Goodacre, Aisha Aljanahi, Subhiksha Nandakumar, Mike Mikailov, and Arifa S Khan. A
275 reference viral database (RVDB) to enhance bioinformatics analysis of high-throughput sequenc-
276 ing for novel virus detection. *mSphere*, 3(2), 2018. doi: 10.1128/mSphereDirect.00069-18.
- 277 Manfred G Grabherr, Brian J Haas, Moran Yassour, Joshua Z Levin, Dawn A Thompson, Ido
278 Amit, Xian Adiconis, Lin Fan, Raktima Raychowdhury, Qiangdong Zeng, Zehua Chen, Evan
279 Mauceli, Nir Hacohen, Andreas Gnirke, Nicholas Rhind, Federica di Palma, Bruce W Birren,
280 Chad Nusbaum, Kerstin Lindblad-Toh, Nir Friedman, and Aviv Regev. Full-length transcriptome
281 assembly from RNA-seq data without a reference genome. *Nat. Biotechnol.*, 29(7):644–652, 2011.
282 doi: 10.1038/nbt.1883.
- 283 Nathan D Grubaugh, Karthik Gangavarapu, Joshua Quick, Nathaniel L Matteson, Jaqueline Goes
284 De Jesus, Bradley J Main, Amanda L Tan, Lauren M Paul, Doug E Brackney, Saran Grewal, Nikos
285 Gurfield, Koen K A Van Rompay, Sharon Isern, Scott F Michael, Lark L Coffey, Nicholas J Loman,
286 and Kristian G Andersen. An amplicon-based sequencing framework for accurately measuring
287 intrahost virus diversity using PrimalSeq and iVar. *Genome Biol.*, 20(1):8, 2019. doi: 10.1186/
288 s13059-018-1618-7.
- 289 Jiarong Guo, Ben Bolduc, Ahmed A Zayed, Arvind Varsani, Guillermo Dominguez-Huerta, Tom O
290 Delmont, Akbar Adjie Pratama, M Consuelo Gazitúa, Dean Vik, Matthew B Sullivan, and Simon
291 Roux. VirSorter2: a multi-classifier, expert-guided approach to detect diverse DNA and RNA
292 viruses. *Microbiome*, 9(1):37, 2021. doi: 10.1186/s40168-020-00990-y.
- 293 Kazutaka Katoh, Kazuharu Misawa, Kei-Ichi Kuma, and Takashi Miyata. MAFFT: a novel method
294 for rapid multiple sequence alignment based on fast fourier transform. *Nucleic Acids Res.*, 30
295 (14):3059–3066, 2002. doi: 10.1093/nar/gkf436.
- 296 Kristopher Kieft, Zhichao Zhou, and Karthik Anantharaman. VIBRANT: automated recovery,
297 annotation and curation of microbial viruses, and evaluation of viral community function from
298 genomic sequences. *Microbiome*, 8(1):90, 2020. doi: 10.1186/s40168-020-00867-0.
- 299 Kristopher Kieft, Alyssa Adams, Rauf Salamzade, Lindsay Kalan, and Karthik Anantharaman.
300 vRhyme enables binning of viral genomes from metagenomes. *Nucleic Acids Res.*, 50(14):e83,
301 2022. doi: 10.1093/nar/gkac341.
- 302 Gregory M Kurtzer, Vanessa Sochat, and Michael W Bauer. Singularity: Scientific containers for
303 mobility of compute. *PLoS One*, 12(5):e0177459, 2017. doi: 10.1371/journal.pone.0177459.
- 304 Björn E Langer, Andreia Amaral, Marie-Odile Baudement, Franziska Bonath, Mathieu Charles,
305 Praveen Krishna Chitneedi, Emily L Clark, Paolo Di Tommaso, Sarah Djebali, Philip A Ewels,
306 Sonia Eynard, James A Fellows Yates, Daniel Fischer, Evan W Floden, Sylvain Foissac, Gisela
307 Gabernet, Maxime U Garcia, Gareth Gillard, Manu Kumar Gundappa, Cervin Guyomar, Christo-
308 pher Hakkaart, Friederike Hanssen, Peter W Harrison, Matthias Hörtenhuber, Cyril Kurylo,
309 Christa Kühn, Sandrine Lagarrigue, Delphine Lallias, Daniel J Macqueen, Edmund Miller, Júlia
310 Mir-Pedrol, Gabriel Costa Monteiro Moreira, Sven Nahnsen, Harshil Patel, Alexander Peltzer,

- Frederique Pitel, Yulixaxis Ramayo-Caldas, Marcel da Câmara Ribeiro-Dantas, Dominique Rocha, Mazdak Salavati, Alexey Sokolov, Jose Espinosa-Carrasco, Cedric Notredame, and The Nf-Core Community. Empowering bioinformatics communities with nextflow and nf-core. *Genome Biol.*, 26(1):228, 2025. doi: 10.1186/s13059-025-03673-9.
- Ben Langmead, Christopher Wilks, Valentin Antonescu, and Rone Charles. Scaling read aligners to hundreds of threads on general-purpose processors. *Bioinformatics*, 35(3):421–432, 2019. doi: 10.1093/bioinformatics/bty648.
- Jeroen F J Laros. HUMID: reference free FastQ deduplication, 2025.
- Dinghua Li, Ruibang Luo, Chi-Man Liu, Chi-Ming Leung, Hing-Fung Ting, Kunihiko Sadakane, Hiroshi Yamashita, and Tak-Wah Lam. MEGAHIT v1.0: A fast and scalable metagenome assembler driven by advanced methodologies and community practices. *Methods*, 102:3–11, 2016. doi: 10.1016/j.ymeth.2016.02.020.
- ~~Heng Li. Aligning sequence reads, clone sequences and assembly contigs with BWA-MEM. *arXiv q-bio.GN*, 2013.~~
- Heng Li. Minimap2: pairwise alignment for nucleotide sequences. *Bioinformatics*, 34(18):3094–3100, 2018. doi: 10.1093/bioinformatics/bty191.
- Weizhong Li and Adam Godzik. Cd-hit: a fast program for clustering and comparing large sets of protein or nucleotide sequences. *Bioinformatics*, 22(13):1658–1659, 2006. doi: 10.1093/bioinformatics/btl158.
- Dmitry Meleshko, Iman Hajirasouliha, and Anton Korobeynikov. coronaSPAdes: from biosynthetic gene clusters to RNA viral assemblies. *Bioinformatics*, 2021. doi: 10.1093/bioinformatics/btab597.
- Peter Menzel, Kim Lee Ng, and Anders Krogh. Fast and sensitive taxonomic classification for metagenomics with kaiju. *Nat. Commun.*, 7:11257, 2016. doi: 10.1038/ncomms11257.
- D Merkel. Docker: lightweight linux containers for consistent development and deployment. *Linux Journal*, 2014(239):2, 2014. doi: 10.5555/2600239.2600241.
- Stephen Nayfach, Antonio Pedro Camargo, Frederik Schulz, Emiley Eloë-Fadrosch, Simon Roux, and Nikos C Kyrpides. CheckV assesses the quality and completeness of metagenome-assembled viral genomes. *Nat. Biotechnol.*, 39(5):578–585, 2021. doi: 10.1038/s41587-020-00774-7.
- Brian D Ondov, Nicholas H Bergman, and Adam M Phillippy. Interactive metagenomic visualization in a web browser. *BMC Bioinformatics*, 12(1):385, 2011. doi: 10.1186/1471-2105-12-385.
- Brian D Ondov, Gabriel J Starrett, Anna Sappington, Aleksandra Kostic, Sergey Koren, Christopher B Buck, and Adam M Phillippy. Mash screen: high-throughput sequence containment estimation for genome discovery. *Genome Biol.*, 20(1):232, 2019. doi: 10.1186/s13059-019-1841-x.
- Guillermo Rangel-Pineros, Alexandre Almeida, Martin Beracochea, Ekaterina Sakharova, Manja Marz, Alejandro Reyes Muñoz, Martin Hölzer, and Robert D Finn. VIRify: an integrated detection, annotation and taxonomic classification pipeline using virus-specific protein profile hidden markov models. *bioRxiv*, page 2022.08.22.504484, 2022. doi: 10.1101/2022.08.22.504484.

- 349 Torbjørn Rognes, Tomáš Flouri, Ben Nichols, Christopher Quince, and Frédéric Mahé. VSEARCH:
350 a versatile open source tool for metagenomics. *PeerJ*, 4:e2584, 2016. doi: 10.7717/peerj.2584.
- 351 Samuel S Shepard, Sarah Meno, Justin Bahl, Malania M Wilson, John Barnes, and Elizabeth
352 Neuhaus. Viral deep sequencing needs an adaptive approach: IRMA, the iterative refinement
353 meta-assembler. *BMC Genomics*, 17(1):708, 2016. doi: 10.1186/s12864-016-3030-6.
- 354 Tom Smith, Andreas Heger, and Ian Sudbery. UMI-tools: modeling sequencing errors in unique
355 molecular identifiers to improve quantification accuracy. *Genome Res.*, 27(3):491–499, 2017. doi:
356 10.1101/gr.209601.116.
- 357 Martin Steinegger and Johannes Söding. MMseqs2 enables sensitive protein sequence searching for
358 the analysis of massive data sets. *Nat. Biotechnol.*, 35(11):1026–1028, 2017. doi: 10.1038/nbt.3988.
- 359 Md Vasimuddin, Sanchit Misra, Heng Li, and Srinivas Aluru. Efficient architecture-aware acceler-
360 ation of BWA-MEM for multicore systems. In *2019 IEEE International Parallel and Distributed
361 Processing Symposium (IPDPS)*, pages 314–324. IEEE, 2019. doi: 10.1109/IPDPS.2019.00041.
- 362 Derrick E Wood, Jennifer Lu, and Ben Langmead. Improved metagenomic analysis with kraken 2.
363 *Genome Biol.*, 20(1):257, 2019. doi: 10.1186/s13059-019-1891-0.
- 364 Andrzej Zielezinski, Adam Gudyś, Jakub Barylski, Krzysztof Siminski, Piotr Rozwalak, Bas E
365 Dutilh, and Sebastian Deorowicz. Ultrafast and accurate sequence alignment and clustering of
366 viral genomes. *Nat. Methods*, 22(6):1191–1194, 2025. doi: 10.1038/s41592-025-02701-7.