

nf-core/viralmetagenome: A Novel Pipeline for Untargeted Viral Genome Reconstruction

Joon Klaps^{1,*}, Philippe Lemey¹, nf-core community², Liana Eleni Kafetzopoulou¹

¹Rega Institute for Medical Research, Department of Microbiology, Immunology and Transplantation, KU Leuven, Belgium

²A full list of contributors can be found at <https://nf-co.re/community>.

*Correspondence: joon.klaps@kuleuven.be

November 24, 2025

Abstract

Motivation: Eukaryotic viruses present significant challenges for genome reconstruction and variant analysis due to their extensive diversity and potential genome segmentation. While de novo assembly followed by reference database matching and scaffolding is ~~a commonly used approach, the manual execution of this workflow is extremely time-consuming, particularly due to the common, manual execution is labor-intensive, particularly when~~ extensive reference curation ~~required. Here, we is required.~~ We address the critical need for an automated, scalable pipeline ~~that can efficiently handle for efficient~~ viral metagenomic analysis without manual intervention.

Results: We present nf-core/viralmetagenome, a comprehensive viral metagenomic pipeline for untargeted genome reconstruction and variant analysis of eukaryotic DNA and RNA viruses. Viralmetagenome is implemented as a Nextflow workflow that processes short-read metagenomic samples to automatically detect and assemble viral genomes, while also performing variant analysis. The pipeline features automated reference selection, consensus quality control metrics, comprehensive documentation, and seamless integration with containerization technologies, including Docker and Singularity. We demonstrate the utility and accuracy of our approach through validation on both simulated and real datasets, showing robust performance across diverse viral families in metagenomic samples.

Availability: nf-core/viralmetagenome is freely available at <https://github.com/nf-core/viralmetagenome> with comprehensive documentation at <https://nf-co.re/viralmetagenome>

Contact: joon.klaps@kuleuven.be

Supplementary information: Supplementary data are available at <https://github.com/Joon-Klaps/nf-core-viralmetagenome-manuscript> online.

Keywords: viralmetagenome, bioinformatic pipeline, nextflow, viral metagenomics, viral assembly, viral variant analysis

1 Introduction

Reconstructing viral genomes from metagenomic sequencing data presents considerable computational challenges, particularly for viruses exhibiting extensive genetic diversity (Baaijens et al. 2017, Meleshko et al. 2021). This ~~diversity is further compounded by segmented genomes in families like~~ challenge is amplified in segmented viruses such as influenza, rotavirus, and bunyaviruses, where individual segments can evolve under distinct selective pressures and ~~reassort, contributing to a complex landscape for genome reassortment~~ complicates reconstruction. While ~~pipelines are often designed to some~~ pipelines target specific viruses and their subtypes (Shepard et al. 2016), accurate ~~and complete~~ genome reconstruction of samples with unknown references typically requires ~~time-consuming~~ manual curation of contigs and reference matching (de Vries et al. 2021). This ~~manual curation process is time-consuming, making it impractical for~~ makes large-scale metagenome studies or rapid ~~response scenarios that involve emerging viral outbreaks of unknown origin~~ outbreak responses impractical.

To address these limitations, we developed nf-core/viralmetagenome, a comprehensive pipeline specifically designed for untargeted viral genome reconstruction. The pipeline is developed using Nextflow (Di Tommaso et al. 2017) within the nf-core framework (~~Ewels et al. 2020~~)(Ewels et al. 2020, Langer et al. 2020), ensuring reproducibility through containerization with Docker (Merkel 2014) and Singularity (Kurtzer et al. 2017), and enabling portability across computational platforms such as local desktops, high-performance clusters and cloud environments.

2 Pipeline Description

nf-core/viralmetagenome implements an automated workflow ~~performing for~~ *de novo* assembly, reference matching, and iterative consensus refinement ~~for the reconstruction of~~ to reconstruct viral genomes without prior target knowledge. The pipeline consists of five major stages: read preprocessing, metagenomic diversity assessment, contig assembly and scaffolding, iterative consensus refinement with variant analysis, and quality control (Figure 1). ~~While this manuscript highlights key differences between particular tools, unless~~ Unless otherwise specified, the pipeline offers multiple options to accommodate established user workflows and preferences. ~~In depth tool details are available~~ Detailed tool descriptions are provided in Supplementary Table 1.

2.1 Read preprocessing

~~Input reads provided via~~ Input reads listed in sample sheets containing sample names and short-read FASTQ paths are preprocessed using FastQC and ~~have their adapters trimmed with~~ trimmed for adapters using fastp (Chen et al. 2018) (default) or Trimmomatic (Bolger et al. 2014). Fastp ~~is overall faster and has~~ offers automated adapter detection and ~~trimming faster~~ processing (Chen et al. 2018). UMI deduplication is implemented using HUMID (Laros 2025) ~~and once reads are mapped to a reference with~~ for raw reads, or UMI-tools after mapping (Smith et al. 2017). Multiple sequencing runs ~~are~~ can be merged after trimming by specifying merge group identifiers in the sample sheet. ~~Complexity filtering~~ Filtering with bbduk (Bushnell 2022) or PRINSEQ++ (Cantu et al. 2019) removes low-complexity sequences ~~containing repetitive elements where~~, while host reads are




Figure 1: ~~Visual overview of the~~ Overview of the nf-core/viralmetagenome pipeline for untargeted viral genome reconstruction. nf-core/viralmetagenome ~~pipeline for untargeted viral genome reconstruction.~~ nf-core/viralmetagenome processes short-read data through pre-processing, metagenomic diversity assessment, ~~de novo~~ *de novo* assembly with multiple assemblers, scaffolding with automated reference identification and ~~contig~~ *contig*-taxonomy-guided clustering of *contigs*, and iterative consensus refinement through read mapping and variant calling. Quality control metrics, assembly statistics, and coverage data are integrated into interactive MultiQC reports and standardised overview tables for downstream analysis.

removed with Kraken2 (Wood et al. 2019).

2.2 Metagenomic diversity assessment

Taxonomic classification of preprocessed reads is performed using two complementary approaches - Kaiju (Menzel et al. 2016) and Kraken2 (Wood et al. 2019) - to maximise detection sensitivity across diverse viral families. Results from both classifiers are visualised using Krona (Ondov et al. 2011).

2.3 ~~De-novo~~ *De novo* assembly and clustering

The assembly workflow ~~implements-employs~~ multi-assembler approaches followed by clustering and scaffolding ~~of contigs independent of the original assembler~~. *De novo* assembly ~~is-performed-using~~ ~~uses~~ SPAdes (Meleshko et al. 2021) (RNAviral mode), MEGAHIT (Li et al. 2016), and Trinity (Grabherr et al. 2011), capitalizing on distinct ~~algorithmic-strengths-algorithms~~ to maximise genome recovery across diverse viral families and ~~variable~~-read depths.

Reference identification uses BLASTn (Altschul et al. 1990) against the Reference Viral Database (RVDB) (Goodacre et al. 2018), retaining ~~the~~ top five hits to ~~facilitate-identification-of-identify~~ related genomic segments ~~and appropriate reference sequences for contig scaffolding and clustering~~. ~~for scaffolding~~. Clustering is performed in two sequential stages. First, taxonomic pre-clustering groups contigs based on taxonomic classification using Kraken2 (Wood et al. 2019) and Kaiju (Menzel et al. 2016), with optional filtering to focus on specific taxonomic clades for more targeted analyses. Second, nucleotide similarity clustering within taxonomic groups is performed using CD-HIT-EST (Li and Godzik 2006) (~~default~~), VSEARCH (Rognes et al. 2016), MMseqs2 (Steinegger and Söding 2017), vRhyme (Kieft et al. 2022), or Mash (Ondov et al. 2019). All tools are valid options, though performance may vary depending on the dataset (Zielezinski et al. 2025, Steinegger and Söding 2017).

As an optional filtering step of contig clusters ~~,after-following~~ assembly and extension, reads can be mapped to all contigs using BWA-MEM2 (Vasimuddin et al. 2019) (default), ~~BWA-(?), or bowtie2~~ ~~or Bowtie2~~ (Langmead et al. 2019). Clusters are filtered based on the total percentage of reads mapping to all contigs within a cluster, allowing identification and removal of clusters that likely represent assembly artefacts resulting from low read coverage.

For the final scaffolding step, all cluster members are mapped to the cluster representative or centroid using Minimap2 (Li 2018), followed by consensus calling with iVar (Grubaugh et al. 2019) to generate reference-assisted assemblies. ~~Regions with zero coverage depth can optionally be represented by the reference genome to produce a more complete scaffold genome for consensus calling~~.

2.4 Iterative consensus refinement and variant calling

The consensus module supports external reference-based analysis and scaffold refinement. Users can provide ~~a separate reference genome or reference set for each sample with mapping constraints~~;

~~when a reference~~ specific reference genomes; if a set is provided, Mash (Ondov et al. 2019) selects the most similar ~~can be selected using Mash (Ondov et al. 2019).~~

~~Scaffold refinement performs one.~~ The selected reference genome undergoes a single round of refinement to generate the consensus. In contrast, scaffold refinement employs an iterative approach, performing up to 4 iterative cycles (default 2). ~~Each iteration maps to polish the consensus sequence.~~ Both processes map reads using BWA-MEM2 (Vasimuddin et al. 2019), ~~BWA (?), or bowtie2 (Langmead et al. 2019) to the consensus~~ or Bowtie2 (Langmead et al. 2019), followed by variant calling with BCFtools (Danecek et al. 2021) or iVar (Grubaugh et al. 2019). Benchmarking ~~by Bassano et al. (Bassano et al. 2022) showed that BCFtools outperformed (Bassano et al. 2022)~~ indicated that BCFtools outperforms iVar in precision and recall, ~~where while~~ iVar identified more ~~low frequency~~ low-frequency variants. Users ~~are recommended to consider prioritising~~ can prioritise sensitivity or specificity when selecting the variant caller. Final variant calling is performed on the polished consensus to characterise single nucleotide variants.

2.5 Consensus Quality Control

Quality control employs CheckV (Nayfach et al. 2021) ~~for completeness estimates to assess completeness,~~ BLASTn (Altschul et al. 1990) for reference similarity, and MMseqs2 (Steinegger and Söding 2017) against the annotated database Virosaurus (Gleizes et al. 2020). These analyses enable species identification, genomic segment classification, host prediction, and any other additional metadata embedded within the reference databases.

The refinement progression is evaluated through sequence alignment with MAFFT (Katoh et al. 2002), which compares final consensus genomes against *de novo* contigs, intermediate consensus sequences from iterative cycles, and the scaffolding reference. All tool metrics are compiled into an interactive MultiQC report (Ewels et al. 2016). Additionally, key metrics are extracted from the MultiQC report and compiled into standalone overview tables to facilitate downstream analysis across all processed samples.

3 Applications

~~To assess the performance of nf-core/viralmetagenome under challenging scenarios, we simulated coinfection scenarios by mixing paired-end reads from public HIV-1 genomes with varying diversity (80-99%), resulting in 13 samples (See supplementary table 2). Nf-core/viralmetagenome successfully identified coinfections in all mixed samples when genetic similarity was low to moderate ($\leq 96.7\%$).~~

~~We validated nf-core/viralmetagenome~~ We validated performance on real metagenomic samples spanning human and plant pathogens. ~~Here, the pipeline successfully nf-core/viralmetagenome generated high-quality or near-complete genomes across viral families including segmented (Lassa virus, Orthonairovirus, Tomato spotted wilt tospovirus) and non-segmented viruses (SARS-CoV-2, West Nile virus, Potato virus Y, Youcai mosaic virus, and Monkeypox virus).~~

~~Processing; Supplementary methods S1).~~ Processing these 28 samples (supplementary methods) required 412 CPU hours and a maximum of 79GB RAM on ~~an HPC, excluding taxonomic classification~~

~~steps. The automated reference selection offers substantial improvements over manual curation by~~
~~reducing processing time while preserving reconstruction accuracy. Performance correlates strongly~~
~~with reference database comprehensiveness, as consensus genomes tended to be more complete and~~
~~similar to the true consensus sequence when a high-performance cluster.~~ For local desktop use, we
 suggest to use the `-profile local` option which disables the memory intensive Kaiju database in
 favor of a smaller one (viral refseq), reducing maximum RAM usage to approximately 18GB.

To assess performance under challenging scenarios, we conducted two targeted validation exper-
 iments. First, we evaluated the reconstruction of multiple viral genomes in a simulated mixture
 sample (Supplementary Methods S2). The pipeline successfully separated coinfecting strains when
 the genetic divergence between them was reasonably high ($\leq 88.7\%$ ANI; Supplementary Figure 1).
 Second, we investigated the influence of the scaffolding reference using real Lassa virus sequencing
 data (Supplementary Methods S3). We first established a baseline consensus reconstruction by
 providing the unclustered-RVDB as the reference input, and then compared this baseline against
 consensus reconstructions generated using single, increasingly dissimilar scaffolding references (rang-
 ing from 64% to 100% ANI).

Our results showed that sequencing depth dictates the importance of the scaffolding reference ~~was~~
~~closer to the true viral genome. This emphasizes the~~. For the high-coverage S segments, *de*
novovo assembly alone was often sufficient to reconstruct the full genome, so the choice of scaffolding
 reference had minimal impact (Figure 2A-B). In contrast, for the lower-coverage L segments, *de novo*
 assembly often resulted in fragmented contigs. Here, the use of a closely related external reference
 was critical, allowing the pipeline to bridge gaps and produce a significantly more complete and
 accurate consensus genome compared to using *de novo* contigs alone (Figure 2C-D).

These findings emphasise the need to keep databases like RVDB (Goodacre et al. 2018) ~~and~~
~~Virosaurus (Gleizes et al. 2020)~~ up-to-date. Since nf-core/viralmetagenome is primarily designed
 for eukaryotic viruses, ~~bacteriophage analysis requires different approaches and users~~ users interested
 in bacteriophage analysis are encouraged to explore pipelines ~~targeting~~ designed to target phages
 such as VIRify (Rangel-Pineros et al. 2022), VIBRANT (Kieft et al. 2020), VirSorter2 (Guo et al.
 2021).

4 Conclusion

nf-core/viralmetagenome addresses a critical need in viral genomics by providing an automated,
 scalable solution for untargeted viral genome reconstruction. The pipeline successfully automates
 the traditionally time-consuming and manual execution process of viral genome ~~assembly-reconstruc-~~
~~tion~~ from short-read metagenomic data through its integrated workflow of *de novo* contig assembly,
 automated reference selection, clustering algorithms, and iterative refinement strategies.

~~Our validation demonstrates the pipeline's broad applicability across diverse~~ Our validation results
 indicate that the pipeline is suitable for a wide range of eukaryotic viral families, ~~achieving it removes~~
 the burden of extensive manual curation, and enables consistent, high-quality genome reconstruction
~~while ensuring reproducibility and ease of deployment across different computational environments.~~




Figure 2: **Impact of scaffolding reference on consensus reconstruction relative to a U-RVDB baseline.** (A, B) Consensus genome ANI (%) to baseline using external reference sequences vs. *de novo* contigs for L and S segments. (C, D) Correlation between scaffolding reference ANI and resulting consensus ANI to baseline for L and S segments, colored by sample.

with reproducible results across varied computational settings. As viral surveillance and outbreak response increasingly rely on metagenomic sequencing, automated pipelines like nf-core/viralmetagenome will be essential for the timely identification of pathogen strains. The pipeline represents a significant step forward in making viral genome reconstruction accessible to researchers without requiring extensive bioinformatics expertise, facilitating broader adoption of metagenomic approaches in viral research and public health applications.

Acknowledgments

J.K, P.L. and L.E.K acknowledge support from the Research Foundation - Flanders (Fonds voor Wetenschappelijk Onderzoek – Vlaanderen, G005323N and G051322N, 1SH2V24N, 12X9222N). Assistance from AI language models (e.g., ChatGPT) was used to enhance the clarity of the manuscript and to support code development.

Author Contributions

J.K. designed and implemented the pipeline, performed validation analyses, and wrote the manuscript. P.L. and L. E. K. supervised the project and provided critical feedback. The nf-core community contributed to maintaining the pipeline. All authors reviewed and approved the final manuscript.

Conflict of Interest

The authors declare no competing interests.

References

- S F Altschul, W Gish, W Miller, E W Myers, and D J Lipman. Basic local alignment search tool. *J. Mol. Biol.*, 215(3):403–410, 1990. doi: 10.1016/S0022-2836(05)80360-2.
- Jasmijn A Baaijens, Amal Zine El Aabidine, Eric Rivals, and Alexander Schönhuth. De novo assembly of viral quasispecies using overlap graphs. *Genome Res.*, 27(5):835–848, 2017. doi: 10.1101/gr.215038.116.
- Irene Bassano, Vinoy K Ramachandran, Mohammad S Khalifa, Chris J Lilley, Mathew R Brown, Ronny van Aerle, Hubert Denise, William Rowe, Airey George, Edward Cairns, Claudia Wierzbicki, Natalie D Pickwell, Myles Wilson, Matthew Carlile, Nadine Holmes, Alexander Payne, Matthew Loose, Terry A Burke, Steve Paterson, Matthew J Wade, and Jasmine M S Grimley. Evaluation of variant calling algorithms for wastewater-based epidemiology using mixed populations of SARS-CoV-2 variants in synthetic and wastewater samples. *bioRxiv*, 2022. doi: 10.1101/2022.06.06.22275866.
- Anthony M Bolger, Marc Lohse, and Bjoern Usadel. Trimmomatic: a flexible trimmer for illumina sequence data. *Bioinformatics*, 30(15):2114–2120, 2014. doi: 10.1093/bioinformatics/btu170.

- 217 Brian Bushnell. BBMap, 2022.
- 218 Vito Adrian Cantu, Jeffrey Sadural, and Robert Edwards. PRINSEQ++, a multi-threaded tool
219 for fast and efficient quality control and preprocessing of sequencing datasets. *PeerJ*, 2019. doi:
220 10.7287/peerj.preprints.27553v1.
- 221 Shifu Chen, Yanqing Zhou, Yaru Chen, and Jia Gu. fastp: an ultra-fast all-in-one FASTQ prepro-
222 cessor. *Bioinformatics*, 34(17):i884–i890, 2018. doi: 10.1093/bioinformatics/bty560.
- 223 Petr Danecek, James K Bonfield, Jennifer Liddle, John Marshall, Valeriu Ohan, Martin O Pollard,
224 Andrew Whitwham, Thomas Keane, Shane A McCarthy, Robert M Davies, and Heng Li. Twelve
225 years of SAMtools and BCFtools. *Gigascience*, 10(2), 2021. doi: 10.1093/gigascience/giab008.
- 226 Jutte J C de Vries, Julianne R Brown, Nicole Fischer, Igor A Sidorov, Sofia Morfopoulou, Jiabin
227 Huang, Bas B Oude Munnink, Arzu Sayiner, Alihan Bulgurcu, Christophe Rodriguez, Guillaume
228 Gricourt, Els Keyaerts, Leen Beller, Claudia Bachofen, Jakub Kubacki, Cordey Samuel, Laubscher
229 Florian, Schmitz Dennis, Martin Beer, Dirk Hoepfer, Michael Huber, Verena Kufner, Maryam
230 Zaheri, Aitana Lebrand, Anna Papa, Sander van Boheemen, Aloys C M Kroes, Judith Breuer,
231 F Xavier Lopez-Labrador, and Eric C J Claas. Benchmark of thirteen bioinformatic pipelines for
232 metagenomic virus diagnostics using datasets from clinical samples. *J. Clin. Virol.*, 141:104908,
233 2021. doi: 10.1016/j.jcv.2021.104908.
- 234 Paolo Di Tommaso, Maria Chatzou, Evan W Floden, Pablo Prieto Barja, Emilio Palumbo, and
235 Cedric Notredame. Nextflow enables reproducible computational workflows. *Nat. Biotechnol.*, 35
236 (4):316–319, 2017. doi: 10.1038/nbt.3820.
- 237 Philip Ewels, Måns Magnusson, Sverker Lundin, and Max Käller. MultiQC: summarize analysis
238 results for multiple tools and samples in a single report. *Bioinformatics*, 32(19):3047–3048, 2016.
239 doi: 10.1093/bioinformatics/btw354.
- 240 Philip A Ewels, Alexander Peltzer, Sven Fillinger, Harshil Patel, Johannes Alneberg, Andreas
241 Wilm, Maxime Ulysse Garcia, Paolo Di Tommaso, and Sven Nahnsen. The nf-core framework
242 for community-curated bioinformatics pipelines. *Nat. Biotechnol.*, 38(3):276–278, 2020. doi:
243 10.1038/s41587-020-0439-x.
- 244 Anne Gleizes, Florian Laubscher, Nicolas Guex, Christian Iseli, Thomas Junier, Samuel Cordey,
245 Jacques Fellay, Ioannis Xenarios, Laurent Kaiser, and Philippe Le Mercier. Virosaurus a reference
246 to explore and capture virus genetic diversity. *Viruses*, 12(11), 2020. doi: 10.3390/v12111248.
- 247 Norman Goodacre, Aisha Aljanahi, Subhiksha Nandakumar, Mike Mikailov, and Arifa S Khan. A
248 reference viral database (RVDB) to enhance bioinformatics analysis of high-throughput sequenc-
249 ing for novel virus detection. *mSphere*, 3(2), 2018. doi: 10.1128/mSphereDirect.00069-18.
- 250 Manfred G Grabherr, Brian J Haas, Moran Yassour, Joshua Z Levin, Dawn A Thompson, Ido
251 Amit, Xian Adiconis, Lin Fan, Raktima Raychowdhury, Qiandong Zeng, Zehua Chen, Evan
252 Mauceli, Nir Hacohen, Andreas Gnirke, Nicholas Rhind, Federica di Palma, Bruce W Birren,
253 Chad Nusbaum, Kerstin Lindblad-Toh, Nir Friedman, and Aviv Regev. Full-length transcriptome
254 assembly from RNA-seq data without a reference genome. *Nat. Biotechnol.*, 29(7):644–652, 2011.
255 doi: 10.1038/nbt.1883.

- Nathan D Grubaugh, Karthik Gangavarapu, Joshua Quick, Nathaniel L Matteson, Jaqueline Goes De Jesus, Bradley J Main, Amanda L Tan, Lauren M Paul, Doug E Brackney, Saran Grewal, Nikos Gurfield, Koen K A Van Rompay, Sharon Isern, Scott F Michael, Lark L Coffey, Nicholas J Loman, and Kristian G Andersen. An amplicon-based sequencing framework for accurately measuring intrahost virus diversity using PrimalSeq and iVar. *Genome Biol.*, 20(1):8, 2019. doi: 10.1186/s13059-018-1618-7.
- Jiarong Guo, Ben Bolduc, Ahmed A Zayed, Arvind Varsani, Guillermo Dominguez-Huerta, Tom O Delmont, Akbar Adjie Pratama, M Consuelo Gazitúa, Dean Vik, Matthew B Sullivan, and Simon Roux. VirSorter2: a multi-classifier, expert-guided approach to detect diverse DNA and RNA viruses. *Microbiome*, 9(1):37, 2021. doi: 10.1186/s40168-020-00990-y.
- Kazutaka Katoh, Kazuharu Misawa, Kei-Ichi Kuma, and Takashi Miyata. MAFFT: a novel method for rapid multiple sequence alignment based on fast fourier transform. *Nucleic Acids Res.*, 30(14):3059–3066, 2002. doi: 10.1093/nar/gkf436.
- Kristopher Kieft, Zhichao Zhou, and Karthik Anantharaman. VIBRANT: automated recovery, annotation and curation of microbial viruses, and evaluation of viral community function from genomic sequences. *Microbiome*, 8(1):90, 2020. doi: 10.1186/s40168-020-00867-0.
- Kristopher Kieft, Alyssa Adams, Rauf Salamzade, Lindsay Kalan, and Karthik Anantharaman. vRhyme enables binning of viral genomes from metagenomes. *Nucleic Acids Res.*, 50(14):e83, 2022. doi: 10.1093/nar/gkac341.
- Gregory M Kurtzer, Vanessa Sochat, and Michael W Bauer. Singularity: Scientific containers for mobility of compute. *PLoS One*, 12(5):e0177459, 2017. doi: 10.1371/journal.pone.0177459.
- Björn E Langer, Andreia Amaral, Marie-Odile Baudement, Franziska Bonath, Mathieu Charles, Praveen Krishna Chitneedi, Emily L Clark, Paolo Di Tommaso, Sarah Djebali, Philip A Ewels, Sonia Eynard, James A Fellows Yates, Daniel Fischer, Evan W Floden, Sylvain Foissac, Gisela Gabernet, Maxime U Garcia, Gareth Gillard, Manu Kumar Gundappa, Cervin Guyomar, Christopher Hakkaart, Friederike Hanssen, Peter W Harrison, Matthias Hörtenhuber, Cyril Kurylo, Christa Kühn, Sandrine Lagarrigue, Delphine Lallias, Daniel J Macqueen, Edmund Miller, Júlia Mir-Pedrol, Gabriel Costa Monteiro Moreira, Sven Nahnsen, Harshil Patel, Alexander Peltzer, Frederique Pitel, Yuliaxis Ramayo-Caldas, Marcel da Câmara Ribeiro-Dantas, Dominique Rocha, Mazdak Salavati, Alexey Sokolov, Jose Espinosa-Carrasco, Cedric Notredame, and The Nf-Core Community. Empowering bioinformatics communities with nextflow and nf-core. *Genome Biol.*, 26(1):228, 2025. doi: 10.1186/s13059-025-03673-9.
- Ben Langmead, Christopher Wilks, Valentin Antonescu, and Rone Charles. Scaling read aligners to hundreds of threads on general-purpose processors. *Bioinformatics*, 35(3):421–432, 2019. doi: 10.1093/bioinformatics/bty648.
- Jeroen F J Laros. HUMID: HUMID: reference free FastQ deduplication, 2025.
- Dinghua Li, Ruibang Luo, Chi-Man Liu, Chi-Ming Leung, Hing-Fung Ting, Kunihiko Sadakane, Hiroshi Yamashita, and Tak-Wah Lam. MEGAHIT v1.0: A fast and scalable metagenome assembler driven by advanced methodologies and community practices. *Methods*, 102:3–11, 2016. doi: 10.1016/j.ymeth.2016.02.020.

- ~~Heng Li. Aligning sequence reads, clone sequences and assembly contigs with BWA-MEM. *arXiv q-bio.GN*, 2013.~~
- Heng Li. Minimap2: pairwise alignment for nucleotide sequences. *Bioinformatics*, 34(18):3094–3100, 2018. doi: 10.1093/bioinformatics/bty191.
- Weizhong Li and Adam Godzik. Cd-hit: a fast program for clustering and comparing large sets of protein or nucleotide sequences. *Bioinformatics*, 22(13):1658–1659, 2006. doi: 10.1093/bioinformatics/btl158.
- Dmitry Meleshko, Iman Hajirasouliha, and Anton Korobeynikov. coronaSPAdes: from biosynthetic gene clusters to RNA viral assemblies. *Bioinformatics*, 2021. doi: 10.1093/bioinformatics/btab597.
- Peter Menzel, Kim Lee Ng, and Anders Krogh. Fast and sensitive taxonomic classification for metagenomics with kaiju. *Nat. Commun.*, 7:11257, 2016. doi: 10.1038/ncomms11257.
- D Merkel. Docker: lightweight linux containers for consistent development and deployment. *Linux Journal*, 2014(239):2, 2014. doi: 10.5555/2600239.2600241.
- Stephen Nayfach, Antonio Pedro Camargo, Frederik Schulz, Emiley Eloë-Fadrosch, Simon Roux, and Nikos C Kyrpides. CheckV assesses the quality and completeness of metagenome-assembled viral genomes. *Nat. Biotechnol.*, 39(5):578–585, 2021. doi: 10.1038/s41587-020-00774-7.
- Brian D Ondov, Nicholas H Bergman, and Adam M Phillippy. Interactive metagenomic visualization in a web browser. *BMC Bioinformatics*, 12(1):385, 2011. doi: 10.1186/1471-2105-12-385.
- Brian D Ondov, Gabriel J Starrett, Anna Sappington, Aleksandra Kostic, Sergey Koren, Christopher B Buck, and Adam M Phillippy. Mash screen: high-throughput sequence containment estimation for genome discovery. *Genome Biol.*, 20(1):232, 2019. doi: 10.1186/s13059-019-1841-x.
- Guillermo Rangel-Pineros, Alexandre Almeida, Martin Beracochea, Ekaterina Sakharova, Manja Marz, Alejandro Reyes Muñoz, Martin Hölzer, and Robert D Finn. VIRify: an integrated detection, annotation and taxonomic classification pipeline using virus-specific protein profile hidden markov models. *bioRxiv*, page 2022.08.22.504484, 2022. doi: 10.1101/2022.08.22.504484.
- Torbjørn Rognes, Tomáš Flouri, Ben Nichols, Christopher Quince, and Frédéric Mahé. VSEARCH: a versatile open source tool for metagenomics. *PeerJ*, 4:e2584, 2016. doi: 10.7717/peerj.2584.
- Samuel S Shepard, Sarah Meno, Justin Bahl, Malania M Wilson, John Barnes, and Elizabeth Neuhaus. Viral deep sequencing needs an adaptive approach: IRMA, the iterative refinement meta-assembler. *BMC Genomics*, 17(1):708, 2016. doi: 10.1186/s12864-016-3030-6.
- Tom Smith, Andreas Heger, and Ian Sudbery. UMI-tools: modeling sequencing errors in unique molecular identifiers to improve quantification accuracy. *Genome Res.*, 27(3):491–499, 2017. doi: 10.1101/gr.209601.116.
- Martin Steinegger and Johannes Söding. MMseqs2 enables sensitive protein sequence searching for the analysis of massive data sets. *Nat. Biotechnol.*, 35(11):1026–1028, 2017. doi: 10.1038/nbt.3988.

- 332 Md Vasimuddin, Sanchit Misra, Heng Li, and Srinivas Aluru. Efficient architecture-aware acceler-
333 ation of BWA-MEM for multicore systems. In *2019 IEEE International Parallel and Distributed*
334 *Processing Symposium (IPDPS)*, pages 314–324. IEEE, 2019. doi: 10.1109/IPDPS.2019.00041.
- 335 Derrick E Wood, Jennifer Lu, and Ben Langmead. Improved metagenomic analysis with kraken 2.
336 *Genome Biol.*, 20(1):257, 2019. doi: 10.1186/s13059-019-1891-0.
- 337 Andrzej Zielezinski, Adam Gudysz, Jakub Barylski, Krzysztof Siminski, Piotr Rozwalak, Bas E
338 Dutilh, and Sebastian Deorowicz. Ultrafast and accurate sequence alignment and clustering of
339 viral genomes. *Nat. Methods*, 22(6):1191–1194, 2025. doi: 10.1038/s41592-025-02701-7.